

Spoiler Court: Interpretable Spoiler Detection via Adversarial Multi-Agent LLM Judgement

Sangmin Song, Hari Kang, Suhyeon Park, Seongwoo Choi,
Sungkyun Im, and SANGO H LEE

GS Neotek

{song0313, harikang, psh7286, sky21c009, renjoy}@gsneotek.com

Abstract

We present **Spoiler Court**, a multi-agent adversarial framework for interpretable spoiler detection in movie reviews. Three LLM agents, acting as the Plaintiff, Defendant, and Judge, engage in a multi-round courtroom debate to identify spoiler spans, challenge them, and deliver a final judgement with an interpretable audit trail grounded in the movie’s plot summary. Unlike prior work that produces only binary labels, Spoiler Court outputs *which sentences* are spoilers, *why* (via plot fact citations), and *how the decision was reached* (via adversarial debate transcripts). To evaluate the effectiveness of our approach, we construct a multi-genre evaluation dataset spanning 4 movies with binary and character-level span annotations. Under controlled conditions, Spoiler Court achieves an average F1 of 0.733 across 4 movies, outperforming other baselines including a dedicated spoiler detection framework.

1 Introduction

Spoiler reviews on movie platforms such as IMDB can severely diminish a viewer’s enjoyment (Loewenstein, 1994). Current spoiler management relies on self-reported tags, which are unreliable, motivating the need for automatic detection. Yet the task is inherently challenging: judging whether a sentence constitutes a spoiler is subjective, and reliable judgments require grounding each claim against the movie’s actual plot rather than the reader’s intuition.

Beyond detection accuracy, *deployment utility* demands more than a binary label. For content moderators and platforms, where spoilers impose both user-experience and commercial costs, it is essential to identify which sentences to redact to improve user experience via selective spoiler hiding. Existing LLM-based methods (Sun et al., 2023; Nishi and Suzuki, 2025) yield only binary classifications, providing no actionable localization.

We address these gaps with **Spoiler Court**, a multi-agent adversarial framework that models spoiler detection as a courtroom trial. Because spoiler judgments are inherently subjective, a single model’s judgment is brittle; we instead let claims be proposed, contested, and adjudicated.

First, an accuser identifies suspected spoiler sentences grounded in plot evidence, a defender challenges each accusation, and the accuser rebuts. An arbiter then delivers a final judgement, producing not only a binary label but also localized spoiler spans with an interpretable audit trail. To rigorously evaluate this approach, we introduce a novel span-level evaluation paradigm and conduct controlled experiments. Our results demonstrate that Spoiler Court consistently outperforms baseline methods. Additionally, our ablation study reveals that the multi-agent adversarial structure itself drives the performance gains.

We summarize our contributions as follows:

- We propose **Spoiler Court**, a novel multi-agent adversarial framework that recasts spoiler detection as a debate, yielding both precise spoiler spans and interpretable, plot-grounded audit trails.
- We introduce a new **span-level evaluation paradigm** featuring a character-level annotation protocol and a corresponding span-IoU metric. This allows us to assess *which* specific text spans constitute spoilers, addressing a critical gap in benchmarks that have so far supported only binary classification.
- We conduct **controlled experiments** in which all methods share the same LLM backbone and the same information (full plot). By isolating the effect of structure, we demonstrate that the multi-agent court framework, rather than a model or information advantage, drives the improvement across our evaluated datasets.




Figure 1: Overview of the **Spoiler Court** pipeline. The Plaintiff () indicts spoiler spans, each linked to a plot fact; the Defendant () challenges each with a strategy from a fixed catalogue (S1-S6); the Plaintiff () rebuts with additional plot evidence; and the Judge () rules each span **UPHELD** or **DISMISSED**, labeling the review a spoiler if any span is upheld.

2 Related Work

Spoiler Detection. Spoiler detection traditionally relied on heavy supervised training (Wan et al., 2019; Zeng et al., 2024; Zhang et al., 2025), which requires extensive annotated data and complex feature engineering. To overcome this data bottleneck, recent approaches such as CAST (Nishi and Suzuki, 2025) explored the usage of LLMs, demonstrating that zero- or few-shot prompting can effectively bypass the need for heavy training. However, these LLM-based methods still constrain the task to simple review-level *binary* classification. They fail to localize specific spoiler spans or justify their predictions with actual plot evidence. In contrast, our framework moves beyond binary labels to directly extract exact spoiler spans and provide retrieval-grounded explanations, achieving faithful explainability without any task-specific training (Lyu et al., 2024; Madsen et al., 2024).

Multi-Agent Reasoning. Multi-agent debate enhances LLM reasoning and factuality by allowing agents to exchange and critique answers over multiple rounds (Liang et al., 2024; Chan et al., 2024). However, recent studies suggest that simply increasing the number of agents is insufficient. Instead, it is claimed that a single strongly prompted

agent can often match generic multi-agent discussions, implying that true improvements stem from structured, adversarial interactions (Wang et al., 2024). For instance, assigning opposing roles to agents with an impartial judge yields more truthful outcomes (Khan et al., 2024), and explicit prosecution-defense-judge specializations have proven highly effective in adversarial domains like legal reasoning (Chen et al., 2025). Motivated by these successes, we extend this structured adversarial paradigm to spoiler detection, effectively addressing the domain’s inherent subjectivity.

3 Method

Our framework, **Spoiler Court**, recasts spoiler detection as a simulated courtroom trial. We employ three specialized agents: the Plaintiff, the Defendant, and the Judge. All agents are powered by a single LLM backbone (Claude Sonnet 4.6 (Anthropic, 2026)), but they are differentiated by their distinct role-playing prompts and asymmetric access to information. The pipeline proceeds through three main stages.

3.1 Case Opening

The trial begins when a user review is submitted to the court, alongside a full plot summary compiled

Method	Hulk			The Shawshank Redemption			Mean Girls			Blood Diamond			Average	
	P	R	F1	P	R	F1	P	R	F1	P	R	F1	F1	IoU
<i>Single-agent methods</i>														
Zero-shot CoT	1.000	0.478	0.647	0.870	0.513	0.645	0.941	0.400	0.561	0.864	0.442	0.585	0.610	—
Few-shot CoT	1.000	0.239	0.386	0.824	0.359	0.500	0.917	0.275	0.423	0.938	0.349	0.509	0.455	—
<i>Spoiler-specific frameworks</i>														
CAST	0.950	0.391	0.554	0.870	0.513	0.645	0.864	0.475	0.613	0.857	0.419	0.563	0.594	—
<i>Multi-agent methods</i>														
MA-Debate	0.958	0.500	0.657	0.864	0.487	0.623	0.850	0.425	0.567	0.895	0.395	0.548	0.599	—
MAD	1.000	0.391	0.563	0.857	0.462	0.600	0.929	0.325	0.482	0.875	0.488	0.627	0.568	—
<i>Spoiler Court (ours)</i>														
Round 1	0.900	<u>0.587</u>	<u>0.711</u>	<u>0.875</u>	<u>0.539</u>	<u>0.667</u>	0.880	0.550	0.677	0.865	<u>0.744</u>	<u>0.800</u>	<u>0.714</u>	0.330
Round 2	0.900	<u>0.587</u>	<u>0.711</u>	0.870	0.513	0.645	0.885	0.575	0.697	0.838	0.721	0.775	<u>0.707</u>	<u>0.370</u>
Round 3	0.903	0.609	0.727	0.880	0.564	0.688	0.920	0.575	0.708	0.872	0.791	0.829	0.738	0.380

Table 1: Performance across 4 movies. Reviews sampled from the IMDB Spoiler Dataset; ground truth is our manual binary + span annotations. All methods use Claude Sonnet 4.6 with full plot. Avg F1 = mean F1 across 4 movies. IoU = avg. char-level span IoU on SPOILER reviews. Baselines produce no spans. **Bold** = best, underlined = 2nd.

from Wikipedia and official trailers. The Plaintiff agent initiates the process by acting as the accuser. It carefully reads the review to identify potential spoiler spans. This indictment specifies the exact text span, the matching plot fact, and reasoning that the detail is absent from the movie’s introduction. If the Plaintiff finds no potential spoilers, the pipeline exits early, and the case is immediately dismissed as a non-spoiler.

3.2 Multi-Agent Debate

During this stage, both the plaintiff and the defendant must articulate arguments that align with their respective objectives. To accurately simulate the perspective of an unspoiled reader, the Defendant agent is deliberately restricted from accessing the full plot summary; it receives only the indicted sentences and the movie’s general introduction. The Defendant counters each accusation by selecting from a predefined set of defense strategies:

- **[S1] Pre-Known:** The detail is already public knowledge from marketing or trailers.
- **[S2] Cliché:** The span describes a generic genre trope.
- **[S3] Subjective:** The statement is a personal viewpoint rather than an objective plot fact.
- **[S4] Misaligned:** The span is incorrectly matched with the plot fact.
- **[S5] Vague:** The sentence is too ambiguous to genuinely ruin the plot.
- **[S6] Concede:** The spoiler is undeniable, leading the Defendant to rationally concede without forcing a weak defense.

Following the defense, the Judge evaluates each span and rules it as upheld or dismissed. Dismissed spans then enter the next round, where the Defendant re-defends and the Plaintiff counters with a targeted rebuttal. This iterative process forces both sides to refine their arguments, effectively filtering out weak accusations while preserving contested ones.

3.3 Final Judgement

At the end of each round, the Judge reviews the original review, the full plot summary, and the debate transcript accumulated so far, then rules each contested span as upheld or dismissed. If at least one span is upheld, the entire review is assigned a binary spoiler label.

A key advantage of this framework is its inherent interpretability. The system does not rely on post-hoc explanation methods; rather, the debate transcript itself serves as the output. By providing explicit plot citations, categorized defense strategies, and a clear judicial rationale, Spoiler Court generates a transparent, auditable trail that directly supports real-world content moderation workflows. Full prompts for all agents are provided in Appendix F.

4 Experiments

4.1 Experimental Design

Dataset We sample 100 reviews per movie from the IMDB Spoiler Dataset (Misra, 2022) across four diverse genres. However, researchers reported that the original self-reported tags are unreliable: over 25% of spoiler-containing reviews are left

Variant	Hulk			The Shawshank Redemption			Mean Girls			Blood Diamond			IoU
	P	R	F1	P	R	F1	P	R	F1	P	R	F1	
Spoiler Court	0.903	0.609	0.727	0.875	0.539	0.667	0.920	0.575	0.708	0.872	0.791	0.829	0.380
w/o Plot Grounding	0.895	0.370	0.523	0.857	0.462	0.600	0.867	0.650	0.743	0.963	0.605	0.743	0.310
w/o Court Simulation	1.000	0.370	0.540	0.850	0.436	0.576	1.000	0.375	0.545	0.880	0.512	0.647	0.249

Table 2: Ablation study. “w/o Plot Grounding” uses only the movie introduction (no full plot). “w/o Court Simulation” uses a single LLM call (no multi-agent structure). Both use the same backbone (Claude Sonnet 4.6).

untagged (Misra, 2022), user tagging criteria are highly inconsistent (Wan et al., 2019), and 36% of labels shift upon manual review (Nishi and Suzuki, 2025). Furthermore, existing datasets only provide document-level binary labels, which cannot support our novel span-level evaluation. To resolve this label noise and establish a localized ground truth, we re-annotate all sampled reviews from scratch. Full details are provided in Appendix A.

Baselines. To ensure a fair comparison, all evaluated methods share the same LLM backbone (Claude Sonnet 4.6) and identical information (the full plot summary). As shown in Table 1, we categorize our baselines into three groups: single-agent methods (**Zero-shot CoT** (Kojima et al., 2022) and **Few-shot CoT**), spoiler-specific frameworks (**CAST** (Nishi and Suzuki, 2025)), and general-purpose multi-agent methods (**MA-Debate** (Du et al., 2023) and **MAD** (Liang et al., 2024)). Further implementation details for each baseline are provided in Appendix C.

Evaluation Metrics. For binary spoiler classification, we report standard Precision, Recall, and F1 scores. To assess localization accuracy, we compute the character-level IoU specifically on reviews labeled as SPOILER.

4.2 Main Results

Table 1 shows that our Spoiler Court consistently achieves the highest F1 score across all movies, averaging 0.733 at Round 3. Multi-round debate improves performance (0.719 $\rightarrow$ 0.733), as the Plaintiff’s targeted rebuttals effectively counter weak defenses. In contrast, single-call baselines struggle with recall, missing over half of the actual spoilers. As our strategy analysis reveals (Appendix D), spoiler detection is heavily fraught with ambiguity (38% ‘Vague’ defenses). Faced with this, single agents default to a conservative “when in doubt, do not flag” strategy and fail to leverage long plot summaries effectively.

Notably, generic multi-agent approaches (Du

et al., 2023; Liang et al., 2024) also fall short, with average F1 score below 0.60. Lacking localized grounding, their agents debate broadly at the document level. By combining role-specialized adversarial structure with span-level plot grounding.

Our ablation study (Table 2) confirms the necessity of each module. First, removing the plot summary causes sharp performance drops in plot-heavy movies like *Hulk* (0.727 $\rightarrow$ 0.523), demonstrating that introduction-level information alone is insufficient. Second, eliminating the multi-agent court simulation drops the average F1 to 0.577. Without the adversarial framework, single-agent configurations revert to near-perfect precision but suffer from extremely low recall (average 0.42), underscoring that our courtroom structure is essential for overcoming the inherent conservatism of single-agent models.

5 Conclusion

In this work, we introduced **Spoiler Court**, a multi-agent adversarial framework designed to tackle the inherent ambiguity and subjectivity of spoiler detection. Because the boundaries of what constitutes a spoiler are often highly debatable, relying on a single model’s verdict proves brittle. As our experiments showed, single-agent methods tend to adopt a conservative approach, yielding high precision but severely limiting recall.

Our findings demonstrate that structuring this task as a courtroom trial effectively overcomes this limitation. By distributing the cognitive load across specialized roles for an accuser, a defender, and an arbiter, we introduce an adversarial pressure that allows claims to be rigorously contested. This structure enables Spoiler Court to successfully navigate the gray areas of spoiler judgments. Ultimately, this framework not only improves overall detection performance but also provides precise, localized spoiler spans backed by interpretable audit trails.

Limitations

A primary limitation of our study is the potential for data contamination, as the underlying LLM may have memorized plot details of the classic movies in our dataset. To properly evaluate true reasoning over parametric memory, future testing must be conducted on unpolluted data, such as newly released or obscure copyright-free films. Furthermore, incorporating a more contemporary dataset would yield a more accurate evaluation reflecting modern review-writing styles.

Ethics Statement

Our reviews come from the publicly available IMDB Spoiler Dataset (Misra, 2022) and may contain offensive user-generated language. Annotations were produced by members of our research group who consented to their research use; the protocol is detailed in Appendix A.

References

- Anthropic. 2026. Introducing claude sonnet 4.6. <https://www.anthropic.com/news/claude-sonnet-4-6>. Accessed: 2026-06-25.
- Chi-Min Chan, Weize Chen, Yusheng Su, Jianxuan Yu, Wei Xue, Shanghang Zhang, Jie Fu, and Zhiyuan Liu. 2024. ChatEval: Towards better LLM-based evaluators through multi-agent debate. In *The Twelfth International Conference on Learning Representations (ICLR)*.
- Guhong Chen, Liyang Fan, Zihan Gong, Nan Xie, Zixuan Li, Ziqiang Liu, Chengming Li, Qiang Qu, Benyou Wang, Xiang Hu, and Shiwen Huang. 2025. AgentCourt: Simulating court with adversarial evolvable lawyer agents. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 5850–5865.
- Yilun Du, Shuang Li, Antonio Torralba, Joshua B Tenenbaum, and Igor Mordatch. 2023. Improving factuality and reasoning in language models through multi-agent debate. In *International Conference on Machine Learning*.
- Akbir Khan, John Hughes, Dan Valentine, Laura Ruis, Kshitij Sachan, Ansh Radhakrishnan, Edward Grefenstette, Samuel R. Bowman, Tim Rocktäschel, and Ethan Perez. 2024. Debating with more persuasive LLMs leads to more truthful answers. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*.
- Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. 2022. Large language models are zero-shot reasoners. In *Advances in Neural Information Processing Systems*.
- Tian Liang, Zhiwei He, Wenxiang Jiao, Xing Wang, Yan Wang, Rui Wang, Yujiu Yang, Shuming Shi, and Zhaopeng Tu. 2024. Encouraging divergent thinking in large language models through multi-agent debate. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 17889–17904.
- George Loewenstein. 1994. The psychology of curiosity: A review and reinterpretation. *Psychological Bulletin*, 116(1):75–98.
- Qing Lyu, Marianna Apidianaki, and Chris Callison-Burch. 2024. Towards faithful model explanation in NLP: A survey. *Computational Linguistics*, 50(2):657–723.
- Andreas Madsen, Sarath Chandar, and Siva Reddy. 2024. Are self-explanations from large language models faithful? In *Findings of the Association for Computational Linguistics: ACL 2024*.
- Rishabh Misra. 2022. IMDB spoiler dataset. *arXiv preprint arXiv:2212.06034*.
- Genta Nishi and Einoshin Suzuki. 2025. Enhancing the performance of spoiler review detection by a LLM with hints. In *Proceedings of the First Workshop on Comparative Performance Evaluation: From Rules to Language Models (R2LM) associated with RANLP 2025*, pages 97–112, Varna, Bulgaria.
- Xiaofei Sun, Xiaoya Li, Jiwei Li, Fei Wu, Shangwei Guo, Tianwei Zhang, and Guoyin Wang. 2023. Text classification via large language models. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 8990–9005.
- Mengting Wan, Rishabh Misra, Ndapa Nakashole, and Julian McAuley. 2019. Fine-grained spoiler detection from large-scale review corpora. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 2605–2610.
- Qineng Wang, Zihao Wang, Ying Su, Hanghang Tong, and Yangqiu Song. 2024. Rethinking the bounds of LLM reasoning: Are multi-agent discussions the key? In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6106–6131, Bangkok, Thailand. Association for Computational Linguistics.
- Zinan Zeng, Sen Ye, Zijian Cai, Heng Wang, Yuhao Liu, Haokai Zhang, and Minnan Luo. 2024. MMoE: Robust spoiler detection with multi-modal information and domain-aware mixture-of-experts. In *Findings of the Association for Computational Linguistics: ACL 2024*.
- Haokai Zhang, Shengtao Zhang, Zijian Cai, Heng Wang, Ruixuan Zhu, Zinan Zeng, and Minnan Luo. 2025. Unveiling the hidden: Movie genre and user bias in spoiler detection. In *Machine Learning and Knowledge Discovery in Databases. Research Track: European Conference, ECML PKDD 2025*.

A Dataset Setup and Annotation Process

To construct our evaluation benchmark, we selected four movies representing distinct genres: *Blood Diamond* (action/drama), *Mean Girls* (comedy), *Hulk* (sci-fi/action), and *The Shawshank Redemption* (drama). For each movie, we sampled 100 reviews from the IMDB Spoiler Dataset, stratifying them by their original self-reported labels and filtering by length (80 to 2000 characters) to ensure a balanced representation.

To establish a high-quality, span-level ground truth, volunteer annotators from authors’ research group who are interested in the movie evaluated the sampled reviews. They were tasked with providing both a document-level binary label and extracting precise character-level spoiler spans. Spoilers were strictly defined as concrete plot facts documented in the full Wikipedia summary or official trailers, excluding details predictable from the movie’s general introduction. Annotators were also permitted to flag ambiguous reviews where the spoiler judgment was not clear-cut.

Following this independent phase, we computed Cohen’s Kappa (κ) for each movie to measure inter-annotator agreement on the binary spoiler labels. As shown in Table 3, all movies show moderate to substantial agreement, with an average κ of 0.677. To finalize the benchmark, the annotators conducted a reconciliation session where reviews with conflicting binary labels, mismatched spans, or ambiguity flags were thoroughly discussed. A senior annotator adjudicated all remaining disagreements. Only the mutually agreed-upon labels and spans were included in the final dataset.

Movie	Cohen’s κ	Agreement (%)
Hulk	0.584	80.0
The Shawshank Redemption	0.852	93.0
Mean Girls	0.564	79.0
Blood Diamond	0.709	86.0
Average	0.677	84.5

Table 3: Inter-annotator agreement (Cohen’s Kappa) on binary spoiler labels across four movies, measured before the reconciliation phase.

B Annotation Guidelines and Interface

B.1 Annotation Guidelines

Annotators were instructed to label a sentence as a spoiler if it reveals a plot fact documented in the Wikipedia plot summary that is not predictable from the movie’s introduction or official trailer-level synopsis. Specifically, annotators followed these criteria:

1. **Binary label:** Assign SPOILER if the review contains at least one sentence that reveals a significant plot point; otherwise assign NOT_SPOILER.
2. **Span extraction:** For each spoiler sentence, mark the exact character-level span boundaries. Spans must align with sentence boundaries and must not split or merge sentences.
3. **What counts as a spoiler:** The identity, motivation, or fate of major characters; the resolution of central conflicts; twists, surprises, or reveals after the first act; deaths, betrayals, transformations, or hidden relationships.
4. **What does NOT count:** The genre or premise stated in the introduction; generic genre conventions (e.g., “the hero wins”); subjective impressions (e.g., “I cried at the end”); vague references that cannot be matched to a specific plot fact.
5. **Ambiguous cases:** If the spoiler judgment is not clear-cut, annotators were permitted to flag the review as ambiguous for later discussion during the reconciliation session.

B.2 Annotation Interface

We developed a custom web-based annotation tool to facilitate efficient and consistent labeling. The interface presents each review alongside the movie’s full plot summary, allowing annotators to verify their judgments against the ground-truth narrative. Figure 2 shows a screenshot of the annotation interface.

C Baseline Implementation Details

To evaluate the effectiveness of Spoiler Court, we compare it against various approaches categorized into three groups. All baselines are implemented using Claude Sonnet 4.6 and receive the same information: the full Wikipedia plot summary and

Spoiler Court — Annotator

Annotator ID: e.g. alice

Load batch: Hulk

Choose file

Export annotations

Load from S3

Clear local data

0/100 annotated · Movie: hulk · Annotator: ?

Hulk (2003)

Introduction (Intro) · Baseline — content up to here is NOT a spoiler

Bruce Banner, a genetics researcher with a tragic past, suffers a lab accident that makes him transform into a raging, giant green monster when angered, making him a target of forces seeking to abuse his power.

▼ Full Plot · For verification — only facts absent from the intro are spoilers

Governmental genetics researcher David Banner tries to improve human DNA. His supervisor, Colonel Thaddeus "Thunderbolt" Ross, forbids human experimentation, so David experiments on himself. His wife, Edith, soon gives birth to their son, Bruce Banner. David realizes Bruce inherited his mutant DNA and attempts to find a cure. After discovering his experiments, Ross shuts down David's research; David rigs Desert Base's gamma reactor to explode as revenge. Believing he is dangerous, David tries to kill Bruce but accidentally murders Edith when she gets between them; the trauma makes Bruce suppress his early childhood memories. Ross arrests and sends David to a mental hospital, putting the four-year-old Bruce into foster care. Mrs. Krenzler adopts him, and Bruce assumes the surname, growing up believing his birth parents are dead. As an adult, Bruce becomes a scientist at the Berkeley Lab with his girlfriend and Ross's estranged daughter, Betty Ross. Representing the private research company Attheon, the shady Glenn Talbot becomes interested in the scientists' nanomed research to create regenerating soldiers for the military-industrial complex. David reappears as a janitor in the lab building to infiltrate Bruce's life. The now-general Ross investigates, becoming concerned for Betty's safety around Bruce. Bruce saves a colleague named Harper from an

Review 97 of 100

hulk_real_059

Spoiler definition: A sentence that reveals plot content **not predictable** from the **Introduction (Intro)** on the right. The intro serves as the baseline of "permissible" information (trailer-level synopsis). Use the **Full Plot** only to *verify* whether a claim in the review actually appears in the plot.

→ Content that appears in the plot but is also in the intro is **NOT** a spoiler.

Click each sentence to mark it as a spoiler, then select the overall verdict at the bottom. Your work is auto-saved in the browser. Use Export to create a backup.

[1] This was an interesting film, but messy.

[2] The plot was all over the place, and the effects were good, except of course the Hulk himself.

[3] A little too obvious with the CGI effects.

[4] I did like the comic book feel to the movie though.

[5] Suggested for fans of the Hulk.

[6] Nolte is great as always.

Overall verdict for this review:

☐ SPOILER (contains spoilers)
☐ NOT_SPOILER (no spoilers)
☐ UNCERTAIN (ambiguous)

☐ I have read this review in full and completed my annotation (must be checked to count as done)

Notes for adjudication (optional)...

← Previous

Save & Next →

Jump to unannotated

Figure 2: Screenshot of the custom annotation interface used for span-level spoiler labeling. The left panel displays the movie introduction and full plot summary for verification. The right panel presents the review text sentence by sentence, where annotators click to toggle spoiler labels (highlighted in red), select an overall verdict, and optionally leave adjudication notes.

the movie introduction. Since none of the baseline frameworks were originally designed for spoiler detection, we describe our domain adaptation for each below.

Single-Agent Methods

- **Zero-shot CoT** (Kojima et al., 2022): The model receives the review and full plot summary, then performs step-by-step reasoning before producing a binary spoiler label. No examples or task-specific instructions beyond the spoiler definition are provided.
- **Few-shot CoT**: Uses an 8-shot prompting setup (4 spoiler, 4 non-spoiler), with demonstration examples drawn from movies *not* in the evaluation set (The Sixth Sense, Titanic, Fight Club, The Godfather) to prevent data leakage. Each example includes a short review, reasoning, and the correct label.

Spoiler-Specific Frameworks

- **CAST** (Nishi and Suzuki, 2025): Originally designed with movie introduction text only, we adapt CAST by providing the full plot summary to match our information condition. The pipeline performs clue extraction, spoiler type selection from a predefined taxonomy, and final classification within a single inference call. We follow the “-r” (without reasoning) variant, which the original paper reports as the best-performing configuration.

Multi-Agent Methods

Both multi-agent baselines were originally designed for factuality verification and translation quality tasks, respectively. We adapt them to spoiler detection as follows:

- **MA-Debate** (Du et al., 2023): Originally proposed for math and factuality tasks, we adapt the symmetric debate framework to spoiler detection. Three agents independently receive

319

the review and full plot, then produce an initial SPOILER/NOT_SPOILER judgment with reasoning. In subsequent rounds (up to 3), each agent sees the other agents’ previous responses and may revise its answer. The process terminates upon convergence or after the maximum rounds, with the final label determined by majority vote. This requires up to 9 LLM calls per review (3 agents $\times$ 3 rounds).

- **MAD** (Liang et al., 2024): Originally proposed for translation evaluation, we adapt the adversarial tit-for-tat debate framework to spoiler detection. An Affirmative agent is assigned the position that the review is *not* a spoiler, while a Negative agent argues that it *is* a spoiler. Over 3 rounds, the two agents exchange arguments responding to each other’s claims. A separate Judge agent then reviews the full debate transcript alongside the plot summary and renders a final verdict. Unlike Spoiler Court, MAD agents debate at the document level with fixed positions and do not produce span-level localization or plot fact citations.

D Defense Strategy Distribution

In this section, we present a detailed analysis of the defense strategy distribution to supplement the conclusions drawn in the main text. By examining the tactics employed by the Defendant during the multi-round debates, we gain deeper insights into how the framework navigates the inherent ambiguity of spoiler detection.

As illustrated in Table 4, S5 (Vague) is the most frequently utilized strategy at 38%, indicating that the vagueness of a review statement is the primary axis of contention in these trials. Furthermore, the use of S6 (Concede) at 18% confirms that the Defendant agent behaves rationally; it appropriately concedes when faced with clear plot spoilers rather than forcing weak or illogical defenses.

Code	Strategy	%
S5	Vague	38%
S3	Subjective	22%
S6	Concede	18%
S1	Pre-Known	11%
S2	Cliché	7%
S4	Misaligned	4%

Table 4: Defense strategy distribution across all 4 movies.

E Example Debate Transcript

Table 5 shows a multi-round debate from *The Shawshank Redemption*. When the Judge dismisses a span, the dismissed span enters the next round where the Defendant re-defends and the Plaintiff rebuts. This example demonstrates the full 3-round adversarial exchange.

Span 1 — “When Andy finally breaks free, you don’t see him at the end sipping vodka on the beach; he is working polishing an old boat and restoring a hotel in the ‘hope’ of running a successful business.” Plaintiff (Indictment): The review reveals that Andy successfully escapes and is found at the end in Zihuatanejo working on a boat and hotel — a third-act reveal. <i>Plot fact:</i> “He finds Andy on a beach, and the reunited friends happily embrace.” / “He dreams of living in Zihuatanejo . . . where he can start anew.”
Round 1 — Defendant (S4 — Misaligned): The Plaintiff’s cited plot fact states Andy is found on a beach embracing Red, but the span’s specific claims — polishing a boat and restoring a hotel — do not match the official plot summary at all. The Plaintiff has fabricated a correspondence.
Round 1 — Judge: DISMISSED. The specific details (boat-polishing, hotel restoration) are not documented in the plot. The span cannot be confirmed as revealing a documented plot fact.
Round 2 — Defendant (S4 — Misaligned): The core claim persists — the plot documents only a beach reunion, not boat or hotel work.
Round 2 — Plaintiff (Rebuttal): The S4 defense fixates on peripheral details (boat vs. beach) while ignoring the central disclosure: “Andy finally breaks free” explicitly reveals the escape, and Zihuatanejo is documented in the plot as Andy’s dream destination. The reviewer’s embellishments do not neutralize the core spoiler.
Round 2 — Judge: DISMISSED. The misalignment between the span’s details and the plot remains problematic.
Round 3 — Plaintiff (Rebuttal): The phrase “Andy finally breaks free” alone constitutes a spoiler — it reveals the central third-act twist (the escape) regardless of whether subsequent details about boats match the plot verbatim.
Round 3 — Judge: UPHELD. “Andy finally breaks free” reveals the central escape, a third-act plot fact not derivable from the introduction. The peripheral inaccuracies about boats and hotels do not neutralize the core spoiler disclosure.
Span 2 — “Brooks failed to see it.” Plaintiff (Indictment): The review implies Brooks’s tragic inability to adapt to freedom, referencing his suicide. <i>Plot fact:</i> “Brooks is paroled in 1954 after serving 50 years, but cannot adjust to the outside world and eventually hangs himself.”
Round 1 — Defendant (S5 — Vague): “Brooks failed to see it” is too vague — a reader cannot determine what Brooks failed to see, let alone that it references his suicide.
Round 1 — Judge: DISMISSED. The phrase “it” has no clear referent without prior knowledge of the film.
Round 2 — Plaintiff (Rebuttal): Spoiler assessment must account for surrounding review context, which supplies the referent of “it.” In context, “Brooks failed to see it” directly signals his tragic inability to adapt, documented as his parole and suicide in the plot.
Round 3 — Judge: DISMISSED. The surrounding context only establishes a thematic discussion about hope and happiness, not Brooks’s parole or suicide. The span remains too vague to constitute a spoiler.

Table 5: Multi-round debate from *The Shawshank Redemption*. Span 1 is initially dismissed due to factual misalignment (S4), but the Plaintiff’s Round 3 rebuttal successfully argues that “Andy finally breaks free” reveals the escape regardless of peripheral inaccuracies, leading the Judge to overturn the ruling. Span 2 remains dismissed across all rounds — “Brooks failed to see it” is genuinely too vague to communicate a specific plot fact.

F Prompt Design

We present the full prompts used for each agent in the Spoiler Court framework. All agents receive the same shared spoiler definition, which is included in the Plaintiff prompt (Section F.1); subsequent prompts omit it for brevity. Template variables are denoted with . . .

F.1 Plaintiff Prompt

```
[System]
You are the Plaintiff in a spoiler-detection trial. Your role is to identify which sentences of a movie review are spoilers,
    ↳ and to formally charge each identified sentence with a citation to the official plot.

SPOILER DEFINITION
A "spoiler" is a description in a review that reveals a significant plot point of the movie which (a) is documented in the
    ↳ official plot summary, and (b) is not predictable from the movie's introduction, trailer-level synopsis, or early-
    ↳ development premise.

A significant plot point includes, but is not limited to:
- the identity, motivation, or fate of major characters,
- the resolution of central conflicts,
- twists, surprises, or reveals after the first act,
- deaths, betrayals, transformations, or hidden relationships.

The following are NOT spoilers:
- the genre or premise stated in the introduction,
- generic genre conventions ("the hero wins"),
- subjective impressions ("I cried at the end"),
- vague references that cannot be matched to a specific plot fact.

GRANULARITY
Spoiler spans are sentence-level. A span is one full sentence, delimited by sentence boundaries in the input. Do not split
    ↳ sentences. Do not merge sentences across spans.

OUTPUT DISCIPLINE
Follow the output schema exactly. Do not add commentary outside the schema. Do not generate any text after the final closing
    ↳ tag of the schema.

YOUR TASK
1. Read REVIEW_TEXT sentence by sentence.
2. For each sentence, decide whether it reveals a significant plot point that is documented in MOVIE_PLOT but NOT predictable
    ↳ from MOVIE_INTRO.
3. Wrap each spoiler sentence with <spoiler>...</spoiler> tags.
4. For each tagged sentence, write an indictment that:
    (a) quotes the SINGLE most specifically matching plot fact from MOVIE_PLOT (verbatim or near-verbatim),
    (b) explains why this sentence reveals that fact,
    (c) confirms that the fact is not stated in MOVIE_INTRO.

SPECIAL ATTENTION
- If REVIEW_TEXT contains an explicit spoiler marker (e.g., "spoiler ahead", "*spoilers*", "warning: spoilers"), the sentence(
    ↳ s) following the marker should be scrutinized closely.

CONSTRAINTS
- Tag at sentence boundaries only.
- Do not tag a sentence unless you can cite a corresponding plot fact.
- If REVIEW_TEXT contains no spoilers, return TAGGED_REVIEW unchanged and an empty INDICTMENT block.
- Do not invent plot facts. Cite only from MOVIE_PLOT.

[User]
MOVIE_TITLE: {movie_title}

MOVIE_INTRO:
{movie_intro}

MOVIE_PLOT:
{movie_plot}

REVIEW_TEXT:
{review_text}

OUTPUT SCHEMA (return exactly this structure, nothing else):

TAGGED_REVIEW:
<the review, with <spoiler>...</spoiler> wrapped around spoiler sentences>

INDICTMENT:
[1] SPAN: "<exact tagged sentence>"
    PLOT_FACT: "<verbatim or near-verbatim excerpt from MOVIE_PLOT>"
    REASONING: "<one sentence explaining the link>"
    NOT_IN_INTRO: <YES | NO>
[2] ...
(If no spoilers, write: INDICTMENT: (none))
```

F.2 Defendant Prompt

```
[System]
You are the Defendant in a spoiler-detection trial. The Plaintiff has already identified specific sentences as alleged
    ↳ spoilers and filed an indictment. Your role is to defend EACH alleged span individually.

(Same SPOILER DEFINITION, GRANULARITY, and OUTPUT DISCIPLINE as the Plaintiff prompt.)

YOUR TASK
For every span in the indictment, choose ONE primary defense strategy from the following catalogue and write a short argument
    ↳ (1-2 sentences) supporting it. You must address every span; you cannot skip.
```

DEFENSE STRATEGY CATALOGUE

(S1) PRE-KNOWN
The fact is already publicly known from the movie's introduction, trailers, posters, marketing, or genre conventions
→ widely advertised before release.

(S2) CLICHE
The fact is a generic genre convention or stock trope (e.g., "the hero sacrifices himself" in a superhero film) that
→ does not reveal a specific plot point.

(S3) SUBJECTIVE
The sentence expresses a subjective impression or evaluation rather than a plot fact, and the apparent plot reference
→ is incidental or metaphorical.

(S4) MISALIGNED
The Plaintiff's cited PLOT_FACT does not actually match the SPAN; the Plaintiff has misread the review.

(S5) VAGUE
The SPAN is too vague or ambiguous to communicate a specific plot point to a reader who has not seen the movie.

(S6) CONCEDE
The defense is untenable; the SPAN is a clear spoiler. Use this only when no other strategy applies.

CONSTRAINTS

- You may ONLY see the spans the Plaintiff has tagged. Do not reference or defend any other part of the review.
- Each defense must explicitly invoke ONE of S1-S6 by code.
- Do not invent plot facts.
- If the Plaintiff filed no indictment, return DEFENSE: (none).

[User]
MOVIE_TITLE: {movie_title}

MOVIE_INTRO:
{movie_intro}

MOVIE_PLOT:
{movie_plot}

PLAINTIFF_INDICTMENT:
{indictment_block}

OUTPUT SCHEMA (return exactly this structure, nothing else):

DEFENSE:

[1] SPAN: "<the alleged span, copied verbatim>"
STRATEGY: <S1 | S2 | S3 | S4 | S5 | S6>
ARGUMENT: "<1-2 sentences>"

[2] ...
(If indictment is empty: DEFENSE: (none))

F3 Plaintiff Rebuttal Prompt

[System]
(Same system prompt as the Plaintiff.)

[User]
The Defendant has filed the following defenses against your indictment. For each span where the Defendant argued for DISMISSAL
→ (S1-S5), you must directly counter their specific argument. Do NOT repeat your original reasoning; instead, explain
→ why the Defendant's defense strategy is wrong for THIS span.

MOVIE_TITLE: {movie_title}
MOVIE_PLOT:
{movie_plot}

DEFENDANT DEFENSES TO REBUT:
{defense_block}

For each span, write a targeted rebuttal that addresses the Defendant's specific strategy and argument. Cite concrete evidence
→ from MOVIE_PLOT.

OUTPUT SCHEMA (return exactly this structure, nothing else):

REBUTTAL:

[1] SPAN: "<exact span>"
AGAINST: "<Defendant's strategy code, e.g. S5>"
COUNTER: "<1-2 sentences explaining why the defense fails, with plot evidence>"

[2] ...
(If no rebuttals needed, write: REBUTTAL: (none))

F4 Judge Prompt

[System]
You are the Judge in a spoiler-detection trial. You have heard the Plaintiff's indictment and the Defendant's per-span defense
→ . Your role is to issue a final ruling.

(Same SPOILER DEFINITION, GRANULARITY, and OUTPUT DISCIPLINE as the Plaintiff prompt.)

YOUR TASK

1. Review the original REVIEW_TEXT, MOVIE_INTRO, MOVIE_PLOT, the Plaintiff's tagged review and indictment, and the Defendant's
→ defense.
2. For each indicted span, decide UPHELD or DISMISSED:
 - UPHELD: the span is a spoiler; the Plaintiff's indictment stands.
 - DISMISSED: the Defendant's argument is sufficient; the span is not a spoiler.
3. The review's overall LABEL is SPOILER if at least one span is UPHELD, otherwise NOT_SPOILER.
4. Provide a brief RATIONALE (<= 3 sentences) summarizing the decision.

EVIDENCE-BASED DEFENSE EVALUATION (important)
A Defendant's defense is "sufficient" only when it provides concrete evidence about the SPECIFIC SPAN, not when it merely
→ labels the span with a category. In particular:

- DO NOT accept S5 (VAGUE) when the span names a specific character, event, relationship, death, transformation, or
 ↪ backstory fact documented in MOVIE_PLOT.
- DO NOT accept S3 (SUBJECTIVE) when the span, despite evaluative framing, presupposes or discloses a concrete plot fact.
- DO NOT accept S1 (PRE-KNOWN) merely because an actor was cast publicly. The relevant question is whether the SPECIFIC plot
 ↪ reveal in the span is derivable from MOVIE_INTRO or trailer-level synopsis.
- DO accept these strategies when the Defendant grounds them in concrete evidence.

REVIEWER SELF-MARKERS (hard veto rule)

If REVIEW_TEXT contains an explicit spoiler marker immediately preceding an indicted span (e.g., "spoiler ahead", "*spoilers
 ↪ *", "warning: spoilers"), the marker is the reviewer's own confession that the following content is a spoiler. This
 ↪ is a HARD VETO rule:

- UPHOLD the indicted span unconditionally -- regardless of which strategy the Defendant invokes.
- The Defendant's arguments CANNOT defeat a marker-flagged indictment.

CONSTRAINTS

- Decide each span independently.
- You may UPHOLD a span even if the Plaintiff's reasoning is weak, as long as the span itself reveals a plot fact.
- You may DISMISS a span even if the Defendant's argument is weak, if the span itself does not reveal a plot fact.
- If the indictment is empty, LABEL is NOT_SPOILER and UPHELD_SPANS is empty.

[User]

MOVIE_TITLE: {movie_title}

MOVIE_INTRO:
{movie_intro}

MOVIE_PLOT:
{movie_plot}

ORIGINAL_REVIEW:
{review_text}

PLAINTIFF_TAGGED_REVIEW:
{tagged_review}

PLAINTIFF_INDICTMENT:
{indictment_block}

DEFENDANT_DEFENSE:
{defense_block}

PLAINTIFF_REBUTTAL:
{rebuttal_block}

OUTPUT SCHEMA (return exactly this structure, nothing else):

RULING:

LABEL: <SPOILER | NOT_SPOILER>

UPHELD_SPANS:

- "<exact span text>"
- ...

DISMISSED_SPANS:

- "<exact span text>"
- ...

RATIONALE: "<= 3 sentences>"